


THE MACHINE · ANNOUNCED AUGUST 10, 2026 · SUBJECT TO FINAL AGREEMENTS



THEMATIC RESEARCH · NVDA · AI INFRASTRUCTURE CREDIT

The Fifth Seat

What's Really Happening With NVIDIA's \$500 Billion Compute-Financing Machine — and Who Ends Up Holding the Risk

TON618 CAPITAL RESEARCH · AS OF AUGUST 14, 2026 · NVDA \$225.30 CLOSE, AUGUST 13, 2026 · MARKET CAP ≈\$5.46T

A follow-up to The Losing Bets That Built the World. Balance-sheet figures are read from NVIDIA's Form 10-Q as filed; CDS levels are press-reported prints, not terminal-verified, and are labeled as such throughout.



\$1 The Verdict

13×


THE GROWTH OF NVIDIA'S
PRIVATE-INVESTMENT BOOK,
YEAR OVER YEAR, WHILE IT
ARRANGED \$500B OF OTHER
PEOPLE'S CAPITAL

NVIDIA is not selling its risk. It is recruiting co-investors so it can run a bigger book. The August 10 financing platforms create a seat at the table no prior infrastructure wave has had — the originator, who earns on the risk without holding it — and the end of the pipe, via investment-grade NAIC designations into insurance general accounts, is the American annuity holder, who will be long the two things nobody can price: **GPU residual value and take-or-pay counterparty credit.**

The risk-offload read — that NVIDIA built this machine to get AI-cycle exposure off its own balance sheet before the music stops — fails on the filings. A seller reduces exposure. NVIDIA grew its non-marketable investment book from **\$3.2 billion to \$42.3 billion in four quarters** — \$17.9 billion added in the April quarter alone, with **\$27 billion more already committed** through fiscal 2027 — *before* arranging \$500 billion of third-party capital. The motive is not balance-sheet fear; it is that NVIDIA's binding constraint has migrated. Not fab capacity. Not end demand. **Its customers' cost of capital.** Labs and neoclouds cannot fund a ~\$750-billion-a-year buildout from cash flow, so the supplier is manufacturing the financing channel that keeps its own order book — \$119 billion of supply commitments, \$95 billion of it due within this fiscal year — clearing. Risk transfer is the mechanism. Demand generation is the motive. And the collateral's largest scheduled loss-driver is NVIDIA's own product roadmap.

EXHIBIT 1 · THE BOOK THE "SELLER" GREW

NVIDIA non-marketable equity securities, carrying value at period end, with committed additions (NVDA)



\$17.9B

added to NVIDIA's private-investment book in a single quarter (10-Q, Note 6). The machine's designer is buying the asset class, not exiting it — the balance was \$3.2B a year earlier, and another \$27B is committed.

STATION 0 · THE PRESSURE GAUGE · READ BEFORE THE MACHINE




§2 The Gauge

You may have arrived at this note the way most readers did: wondering what NVIDIA's credit-default swaps are trying to say. Four facts, in order of how much they explain:

1. **The market is nine months old — and NVIDIA only recently gave it something to trade.** NVDA 5-year CDS began actively trading in **November 2025**; there was no meaningful market in insuring NVIDIA's debt for the company's entire prior history. Then on **June 18, 2026, NVIDIA printed its first bond deal since 2021 — \$25.0 billion across seven tranches (2028–2056), its largest ever** — roughly tripling total debt to ~\$38 billion including leases, and handing the CDS market both deliverable paper and \$25 billion of fresh new-issue risk for dealers to hedge. But the raise added the same \$25 billion to the cash pile: pro forma, NVIDIA still holds roughly **\$68 billion of net cash** and earned \$58.3 billion in a single quarter. In solvency terms there remains almost nothing to insure; in market-structure terms, there is now something to trade — and the widening began weeks after the deal priced.
2. **The move is violent but the level is benign.** Roughly 40bps in late June, a record in the low 80s in late July — the largest single-day widening since trading began — and ~80bps into the August 10 announcement. At one point the spread crossed *above* Alphabet's (~69bps): a spread ordering no credit analyst would defend on fundamentals. Even at the record, ~80bps prices roughly a 6–7% cumulative five-year default probability at standard recovery assumptions — high-yield-tourist pricing on a AAA-shaped balance sheet.
3. **Each widening leg has coincided with NVIDIA adding obligations its old balance sheet never carried** — first the \$25 billion funded issuance in June, then the contingent stack: the financing platforms, the ecosystem investments, the lease guarantees. The CDS is repricing NVIDIA's *entanglement with its customers' leverage* — and hedging its newly minted paper — not its standalone solvency.
4. **The dominant flow is proxy hedging.** Funds warehousing unrated neocloud, datacenter, and AI-adjacent private paper cannot buy protection on the actual borrowers, so they short the one liquid, correlated name. NVDA CDS is functioning as the AI complex's credit-hedging instrument — the tell being that Oracle, Alphabet, and even SpaceX CDS printed records in the same tape.

EXHIBIT 2 · THE GAUGE: A NINE-MONTH-OLD INSTRUMENT AT ITS RECORD

NVDA 5-year CDS — four dated press-reported prints, straight segments, nothing interpolated (NVDA)



The market has noticed the same structure this note describes — Michael Burry called the platforms "a sign of desperation" — but the desperation read requires NVIDIA to be *shedding* exposure, and the filings show the opposite (§1). The honest reading: the CDS is not a solvency signal; it is the **price of the missing instrument**. There is currently no way to short AI credit directly, so its price is being discovered in the nearest liquid proxy — and the moment rated compute ABS exists, that flow has a better home. Hence audit point 4 (§7): transfer is real only if NVDA protection tightens *relative to* the levered AI names once the platforms issue.

Three days after the record print, NVIDIA answered — though not in the way the protection buyers were positioned for.

STATION I · INTAKE · ANNOUNCED AUGUST 10



§3 The Machine

On August 10, NVIDIA announced six independent financing platforms with **Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs, and KKR**, to mobilize **over \$500 billion of third-party capital** for AI infrastructure — compute as collateral, capital deployed (per press accounts of the structure) through private offerings and bonds issued by special-purpose entities, all of it "subject to execution of the final agreements." Jensen Huang's pitch is the collateral case verbatim: NVIDIA compute is "*broadly adopted, flexible across models and workloads, fungible and transferable across customers and operators, and continuously improved through CUDA*" — and the release's own copy, a sentence ahead of the quote and not Huang's words, adds the superlatives: the lowest token cost, the highest revenue, and the **longest life**. Apollo's Jim Zelter supplied the asset-class case: modern compute is "*a scarce, mission-critical asset class with compelling investment characteristics*." The supplier asserts the collateral is durable; the originator asserts it is investable; neither will hold it.

Separately — and deliberately not NVIDIA's — the price-discovery venue already exists. **Compute Exchange**, launched January 2025 by DRW's Don Wilson and Suna Said of Woodside AI, runs reserved-capacity auctions across 100-plus providers, offers forward contracts on future capacity, and in July 2026 opened a **secondary market in used H100 and A100 systems** (by the venue's own account, used B200s are not yet even available). The rails and the tape are being built by different hands — because the supplier does not want to own the printed mark on its own hardware's residual value.

One more filing fact frames why the machine had to be built. Three direct customers now account for **21%, 17%, and 16% of NVIDIA's revenue — 54% across three checks** — up from two customers at 16% and 14% a year earlier. A financing channel that broadens the buyer base is not a convenience for this order book; it is a structural necessity.



THE OPERATORS' MANUAL · WHO SITS WHERE

§4 The Seats

The Losing Bets That Built the World established the recurring cast of every infrastructure wave: **builders** lose in aggregate, **suppliers** are paid in cash during the boom, **consolidators** buy the wreckage, **users** keep the surplus. In every prior wave, the financier was one of the four in costume — Insull *was* the builder, financing the grid through his own holding-company pyramid; Morgan was the consolidator, financing rail *after* failure at the reorganization price. The August 10 structure introduces a seat with no historical occupant: the **originator** — six firms whose economics are fees on origination, structuring, and management, with the risk held only in transit. The closest rhyme is not Insull or Morgan but the mortgage originators of 2003–2007 — originate-to-distribute, industrialized — with one improvement (four of the six own or manage the end-holder, so they eat some of their own cooking) and one aggravation (the collateral depreciates on the supplier's own product cadence).

Seven seats at the machine — what each gets, and what each is actually exposed to

SEAT	OCCUPANT	WHAT THEY GET	WHAT THEY'RE ACTUALLY EXPOSED TO
Supplier	NVIDIA	Customers' purchasing power underwritten by the deepest capital pools in the system, without consolidating their debt	Its "longest life / fungible" collateral claims are now load-bearing under \$500B of paper; a P&L already ~23% marks on its own investees; its product cadence is a credit event in the collateral pool
Builders	Neoclouds	Cheaper term capital; a refinancing bid for their debt ladders	Disintermediation. Their moat is privileged access to allocation and capital, rented at a spread; a platform that funds anyone's cluster at a rated spread industrializes the scarcity away (see <i>The Thin Floor</i>)
Builders (split)	Hyperscalers	Little directly — they don't need SPE capital	Their capital advantage compresses; SPE-financed GPUs that must be filled become marginal price-cutters in spot compute if demand disappoints
The Fifth Seat	The six GPs	Origination, structuring, and management fees on \$500B; risk held only in transit	The 2003–07 originate-to-distribute seat — with the difference that four of six own the end-holder, so repudiation of the asset class strikes their AUM franchise
End holder	Insurers → annuity savers	Rated spread product; NAIC-designated ABS for general accounts (Athene, Global Atlantic, Brookfield Wealth Solutions, insurance SMAs)	Not "does AI work." GPU residual value at term, and take-or-pay counterparty credit — the two unpriceable variables that tranching makes look ratable

SEAT	OCCUPANT	WHAT THEY GET	WHAT THEY'RE ACTUALLY EXPOSED TO
The reader	Passive investors	—	The same household is now long this collateral twice: equity via NVDA at ~8% of the S&P 500, credit via the annuity general account. A correlation on no risk report anywhere in the chain
Sell side	Dealers, agencies	A genuine new asset class, rated inventory, and at last a direct instrument for AI credit	The curve they build is the mark-to-market machine (§5)

Seats one through four are as ranked in the prior note. The fifth is new — and it is the seat that, in every prior cycle that had even a partial version of it, won.



THE FLAW IN THE MACHINE · RESIDUALS VS THE ROADMAP

§5 The Contradiction

A liquid compute market is the **precondition** for the asset class and the **detonator** under it. Rating agencies cannot rate compute-backed paper without observable liquidation values — which is precisely what Compute Exchange's forwards and secondary tape provide. But pre-mark, collateral is worth what the model says; post-mark, it is worth what the screen says. ABX did not cause the subprime losses; it made them instantaneous, public, and undeniable. The compute forward curve will do the same for GPU residuals.

And the residuals have a scheduled adversary: **the supplier's own roadmap**. NVIDIA needs H100 and H200 spot rates high to support the residual assumptions in the paper its platforms will originate — and it needs to ship Rubin, which obsoletes them, on an annual cadence it cannot slow without ceding the frontier. NVIDIA's product roadmap is the largest single loss-driver in the collateral pool being financed on NVIDIA's collateral narrative. Early price discovery already leans the wrong way: Compute Exchange reports secondary demand concentrated in **H100 and A100 systems rather than Blackwell** — the residual curve being discovered in real time, sloping against the "longest life" claim. Our depreciation series (*The First Invoice*) documented the useful-life debate; that debate now has a \$500-billion credit structure built on one side of it.

The guarantee question is not hypothetical either. NVIDIA's current filings already disclose a template: beginning in fiscal 2026, the company **guarantees partners' facility-lease obligations in the event of their default, in exchange for warrants** — maximum gross exposure \$3.5 billion, terms of five to seven years, \$712 million escrowed by the partners, the guarantees classified as credit derivatives with fair-value changes running through Other income. This is the exact shape of the question that decides the whole thesis. If residual floors, buyback undertakings, or first-loss support of this shape appear inside the SPE platforms, **the \$500 billion is distribution, not transfer** — the risk rides third-party balance sheets while economically remaining NVIDIA's.



§6 The Return Line

Everything above rests on filings and announcements. This section is conjecture, labeled as such — our best guess at how the machine behaves, stated before the SPE documents exist so it can be scored later.

The clean-agency story is the no-free-lunch story: the terminal holder of the paper earns the spread because they accepted the downside — you do not get the cash flow from the hottest asset class in the world without owning its risk. We think that story is directionally right and wrong in three compounding ways.

The buyer will own the risk but not be paid for it. "Accepting the downside" and "being compensated for the downside" are different claims, and the securitization apparatus exists to open a gap between them. The 2006 AAA CDO buyer also accepted the downside — for 30 basis points over Libor. Tranching, an investment-grade NAIC designation, and a collateral class with five quarters of operating history are how a fat, unpriceable risk is sold at a thin, priceable spread. The yield will be calibrated to the rating; the risk is calibrated to reality. The gap between those two calibrations *is* the fifth seat's margin.

The buyer is short an option the counterparty controls. The feature with no historical precedent. A 2006 mortgage-bond buyer took house-price risk that nobody controlled. The compute-paper buyer takes obsolescence risk whose timing is set by one company — NVIDIA decides when Rubin ships and how hard it deprecates the pool backing the paper — and that same company sponsors the financing channel and sells the collateral. The downside's trigger sits in the seller's product-roadmap meeting. No rated spread prices that correctly, because ratings have no cell for "the sponsor schedules the impairment."

In stress, the risk boomerangs to the center. Structured investment vehicles were off-balance-sheet until December 2007, when Citigroup took them back rather than let the franchise die. Every sponsor here has the same compulsion: Apollo cannot let Athene-held compute paper impair without damaging the machine that funds its entire model, and NVIDIA cannot let a major offtaker default and print a loss onto the collateral curve under \$500 billion of paper — it is *already* backstopping partner leases in exchange for warrants. So our base case is a three-layer ownership: **the annuity holder owns the risk contractually, the sponsors own it reputationally, and NVIDIA owns it economically** — because any default in the pool impairs the collateral narrative its order book runs on. The music-stops scenario is not a clean loss eaten by a distant bondholder; it is the discovery that the risk was never severed from the center. Every prior cycle ended with the same discovery.

Questions we cannot yet answer, listed so they can be checked off as the documents surface: Who holds the **equity and first-loss tranches** — retained by the GPs, sold to yield-hunters, or quietly taken by NVIDIA in the warrants-for-guarantees shape it already uses? First-loss placement is the honest map of who believes what. What **residual-value assumption** clears a rating committee for an asset whose secondary tape already trades off "longest life"? Is there any structure in which NVIDIA can *credibly* decline to support a failing offtaker — and if not, is the \$500 billion economically recourse to NVIDIA regardless of what the documents say? If the paper prices wide, the platforms fail to scale; if it prices tight, the risk is mispriced at \$500 billion of scale — **is there any spread at which this paper is both sellable and fairly priced?** And when the first impairment prints, does the NAIC re-designate the class — and what does a forced insurance-general-account seller do to a collateral curve with a nine-month history?

**\$7 Audit Points**

Five falsifiable checks, pre-registered August 13, 2026 — the first scoring at NVIDIA's Q2 FY27 10-Q, expected in late August 2026.

Registered to the trigger monitor; what each outcome would mean

#	THE CHECK	IF IT FIRES
1	The commitments/guarantees footnote extends beyond the \$3.5B facility-lease program to any SPE-platform support — residual floors, buybacks, first-loss, warrant-compensated credit derivatives tied to the platforms	Distribution - not - transfer confirmed; the risk never left
2	The non - marketable investment book exceeds ~\$55B (vs \$42.3B reported plus \$27B committed)	Leverage - not - exit confirmed — the co-investor read of \$1 holds
3	The first rated compute ABS prints: its residual - value assumption, tranche NAIC designations, and whether NVIDIA appears anywhere in the waterfall	Sizes the pipe and names the bag
4	NVDA CDS tightens <i>relative to</i> Oracle and levered AI-adjacent credit after the platforms issue	Transfer is real; the proxy-hedge flow found its home. NVDA and ORCL staying locked together = nothing moved
5	Direct-customer concentration (21/17/16 = 54%) broadens as SPE-financed buyers appear	The machine is generating new demand rather than funding the same three checks

If audit points 1 fails to fire, 5 broadens, and 4 decouples — the three-part falsification — the transfer is genuine, the annuity holder is being paid roughly fairly, and \$6's boomerang is wrong. We would rather record that outcome than have hedged the claim.

§8 What Is Observed, Computed, and Conjectured

Observed (primary sources, independently verified): every balance-sheet and footnote figure — the non-marketable investment ramp and its \$17.9B single-quarter addition, the \$27B of commitments, the \$30.2B marketable equity book and its lockups, the \$13.4B and \$2.6B unrealized-gain components of Other income, the \$3.5B facility-lease guarantee and its \$712M escrow, the 21/17/16 customer concentration, the \$119B supply and \$30B cloud commitments, net cash and the undrawn \$25B commercial paper program — all from NVIDIA's Form 10-Q for the quarter ended April 26, 2026, as filed. The June 18, 2026 \$25.0B seven-tranche senior notes offering is from NVIDIA's Form 8-K as filed. The August 10 announcement facts and both executive quotes are from the NVIDIA and Blackstone press releases as published. An independent claim-verification pass was run before publication; it corrected three items in our own draft — the June bond deal had been omitted and total debt stated at the stale 10-Q figure, a commercial paper program had been mislabeled a revolver, and B200/H200 had been wrongly included in the used-secondary-market claim — and confirmed every 10-Q-sourced figure exactly, including recomputing the DSO and inventory-day table to the decimal from four separate primary filings. We disclose our own corrections for the same reason we count everyone else's.

Computed (ours, method stated): days sales outstanding and inventory days (formula in §9); the ~23% share of pretax income attributable to investment marks (\$13.4B + \$2.6B against \$69.9B); the ~6–7% cumulative default probability implied by an ~80bps five-year spread (method in §9); pro-forma total debt and net cash after the June issuance (10-Q balance plus the 8-K's \$25.0B of notes, cash up by the same proceeds).

Press-reported (not independently verified): all CDS levels and dates, including the November 2025 trading start, the late-July record prints (accounts disagree between roughly 82 and 84bps — itself a comment on the market's depth), the Alphabet crossover, and the Oracle/Alphabet/SpaceX records; the SPE/bond structure of the platforms (the releases themselves specify almost nothing structurally); Compute Exchange's founding details and its reported H100/A100-concentrated secondary demand; the Burry quotation.

Conjectured (labeled): all of §6. The three-layer ownership claim and the boomerang are stated as speculation and pre-registered against §7's audit points rather than asserted as findings.

§9 Sources & Method

SEC filings. NVIDIA Corp. Form 10-Q for the fiscal quarter ended April 26, 2026, filed May 20, 2026, accession 0001045810-26-000052, retrieved via SEC EDGAR (edgartools), pulled August 13, 2026. NVIDIA Corp. Form 8-K filed June 18, 2026, accession 0001193125-26-275783 (Item 8.01) — the \$25.0 billion senior notes offering in seven tranches maturing 2028–2056, verified tranche-by-tranche; pulled August 14, 2026. Figures read from the notes to the condensed consolidated financial statements: cash equivalents and marketable securities including the lockup disclosures (Note 5); non-marketable securities, the adjustments rollforward, equity-method infrastructure funds, and investment commitments (Note 6); the facility-lease guarantee disclosure; commitments and contingencies including supply, cloud, and other vendor commitments (Note 10); concentration of revenue (direct-customer percentages); and the condensed statements themselves (revenue, pretax income, Other income (expense) net, cash and investments, debt). Prior-year comparatives from the same filing.

Announcement. NVIDIA newsroom and Blackstone press releases, both dated August 10, 2026, retrieved August 13, 2026. Both quotes verbatim from the NVIDIA release. The releases do not themselves describe SPEs, bonds, or tranching; that structural description is attributed to contemporaneous press coverage (Reuters explainer; Benzinga) and is labeled press-reported in §8.

Market data. NVDA closing price \$225.30 and market capitalization $\approx$ \$5.46T as of August 13, 2026 (NASDAQ close, via Financial Modeling Prep). Quarterly income-statement and balance-sheet aggregates used for the DSO/inventory computation from Financial Modeling Prep's standardized statements for the four quarters ended July 2025 through April 2026, pulled August 13, 2026. $DSO = \text{accounts receivable} \div \text{quarterly revenue} \times 91$; $\text{inventory days} = \text{inventory} \div \text{quarterly cost of revenue} \times 91$. S&P 500 weight for NVDA ($\sim$ 8%) as of August 6, 2026, per SPY holdings, carried from the prior note.

CDS. All credit-default-swap levels are press-reported prints (Benzinga, July 2026; Yahoo Finance/Reuters explainer, August 2026; Global Markets Investor, August 2026) — no terminal source was used, and the levels should be treated as approximate. The implied-default arithmetic assumes spread $\approx (1 - R) \times \lambda$ with recovery $R = 40\%$: $\lambda \approx 80\text{bps} \div 0.6 \approx 1.33\%/yr$, cumulative five-year default probability $\approx 1 - e^{(-5\lambda)} \approx 6.4\%$. The chart in §2 connects four dated press prints with straight segments and interpolates nothing between them.

Compute Exchange. compute.exchange (platform descriptions, forward contracts); SiliconANGLE, July 17, 2026 (secondary-market launch, founding details, demand concentration). Press-reported, per §8.

Prior TON618 notes relied upon: The Losing Bets That Built the World (the four-seats framework, index-concentration figures); The Neocloud Fuse (neocloud credit); The Thin Floor (CoreWeave's intermediation economics); The First Invoice (GPU useful-life debate); The AI Capex Payback Clock (hyperscaler free-cash-flow rollover); The Private Credit Reckoning (the insurance general-account channel); The Round Trip (Oracle). Chart sourcing lives in this appendix by design; charts carry no source captions.

Assumptions and estimates. The $\sim$ \$750B/yr hyperscaler datacenter-spend figure is a press-reported industry estimate, not a computation of ours. The 2003–07 originate-to-distribute comparison and the ABX and SIV histories are interpretive frames, sourced from the public record of those episodes. Historical spread references (the 30bps AAA CDO spread) are order-of-magnitude period figures.

§10 Disclosures

Information only. This material is published by TON618 Capital for information purposes only. Nothing here is an offer to sell or a solicitation of an offer to buy any security, fund interest, or digital asset, and nothing here is personalized investment advice or a recommendation regarding any instrument.

Publisher's exclusion. All research is published solely as general, impersonal information of regular circulation. It is not tailored to the objectives or circumstances of any individual and is not issued in connection with compensation from any client. The Fund has no clients and distributes all research free of charge. On that basis it publishes in reliance on the publisher's exclusion from the definition of "investment adviser" under the Investment Advisers Act of 1940 (§202(a)(11)(D); cf. *Lowe v. SEC*, 472 U.S. 181 (1985)).

Registration & conflicts. TON618 Capital is not registered as an investment adviser or broker-dealer in any capacity. The Fund is a Bitcoin fund and may hold or transact in the securities or digital assets it discusses; material conflicts are disclosed where they exist. This note discusses, among others, NVIDIA, Oracle, CoreWeave, Alphabet, Microsoft, Amazon, Meta, Broadcom, Apollo Global Management, BlackRock, Blackstone, Brookfield, Goldman Sachs, and KKR; as of the date of this note the Fund holds **no position, long or short, in any single security named herein**. It maintains a strategic Bitcoin allocation, and to the extent this note bears on liquidity conditions, credit formation, and capital-cycle dynamics that inform that allocation, the reader should weigh that as a material interest. The Fund receives no compensation from any party in connection with its research.

Use of AI. Artificial intelligence is used in the creation of this research. **A material conflict is disclosed here: this note analyzes the financing of AI compute infrastructure, and the AI tooling used to produce it is made by a company that is among the larger buyers of the compute whose financing this note discusses.** All methodology and data integrity are reviewed and approved before publication by TON618 Capital's Chief Investment Officer, **Keyth Beck**; errors may nonetheless occur, and readers should verify independently.

CFA. This report was prepared to align with CFA Institute analytical standards (methodology only). CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute. That reference describes the analytical framework applied; it does not imply the report was prepared, reviewed, or authored by a CFA charterholder, and the report is not issued, reviewed, endorsed, certified, or approved by — nor affiliated with — CFA Institute.

Risk & feedback. Past performance is not indicative of future results. Digital assets and equities are volatile and may result in total loss of capital. Corrections and feedback are welcome — please direct them to CIO **Keyth Beck** at keyth@ton618capital.com.

Version 1.0 · published August 14, 2026 · a follow-up to The Losing Bets That Built the World (August 12, 2026) · independently claim-verified before publication.