

EQUITY VALUATION · ENTERPRISE SOFTWARE & AI INFRASTRUCTURE

Oracle — The Round Trip

After a 64% Fall From Its AI-Hype Peak, Oracle No Longer Prices the Nameplate. It Prices the Base Case — Which Leaves No Margin of Safety on a Margin the Company Will Not Disclose.

TON618 CAPITAL RESEARCH · AS OF JULY 16, 2026 ·

A valuation, not a trade. Fair value is an estimate with wide bands, not a price target. Data: Oracle's FY2026 10-K and prior filings (SEC EDGAR, CIK 0001341439), the Q4 FY2026 release and call, S&P Global, and market data as of the July 16, 2026 close (\$124.21, Schwab). Fact ("Data shows...") is separated from opinion ("We believe..."). Figures are DISCLOSED (in a filing), REPORTED (press), CONSENSUS (sell-side), or ESTIMATED (ours) — labelled where it matters. This note is the first of two single-name tests designed by our thematic note *The Neocloud Fuse* (July 16, 2026); the companion, *The Depreciation Question*, supplies the useful-life analysis §9 leans on. Independently fact-checked before release; the Source & Verification Note documents that pass and what we could not verify.

§0 Verdict

FV ≈ \$114

VS \$124.21 — THREE METHODS
SPAN \$104–120; ~8% DOWNSIDE,
RISK DISTRIBUTION SKEWED
DOWN

Oracle is a high-quality software annuity with a debt-funded AI-infrastructure option bolted on, and the market now prices the option at its base case — no better, no worse. Our sum-of-the-parts puts legacy Oracle at roughly **\$72 a share** on its own; at \$124 you are paying about **\$52 for the OCI/AI bet**, which is very nearly our base-case value for it. **Three independent methods triangulate the whole:** sum-of-the-parts **\$119**, a consolidated DCF **\$104** (it charges the near-term cash burn in full), peer P/E multiples **\$120** — a central estimate near **\$114**, about 8% below the price. After a **64% fall** from its September-2025 peak, Oracle is no longer priced for the nameplate backlog — but it is priced for a base case that assumes OCI **gross** economics running roughly **2.5–3× the ~14% the company has let slip**. The point estimate is close to fair; the distribution is not symmetric — the bull case needs a margin expansion the disclosed trend contradicts, while the bear case is well supported. **You are paying full value for an outcome that must go right.**

SCENARIO	WEIGHT	OCI FY30 REV	OCI MATURE MARGIN	IMPLIED VALUE	VS \$124.21
Bull — Catz ramp broadly hits, margin nears the 30–40% aspiration	22%	\$130B	32%	\$176	+41%
Base — ramp slips ~30%, margin settles below aspiration	43%	\$95B	22%	\$122	–2%
Bear — thin margins persist, capex treadmill, counterparty strain	35%	\$60B	10%	\$80	–36%
Sum-of-the-parts fair value (Method 1 of 3)				≈ \$119	–4%

*Two cross-checks (§7) bracket this: a consolidated DCF — which fully charges the ~\$63B of near-term cash burn the SOTP's exit multiple glosses — lands at **~\$104**, and peer P/E multiples at **~\$120**. We carry **~\$114**, the mean of the three, as the central estimate. (The multiples that make Oracle look cheaper still — EV/EBITDA, EV/Revenue — are cheap precisely because they exclude the capex; §7 explains why they get no weight.)*

What this note is, and is not. It is a valuation of the equity, built to answer the specific question our thematic note handed it: *has the equity priced the risk that the AI backlog converts to far less profit than its headline implies, or is it capitalizing the nameplate as though the contracts were floors?* For Oracle the answer is unusually clean, and it is not the answer the thematic note's prior expected. The equity **already corrected** — hard — and now sits at a defensible base-case value. It is not obviously cheap and not obviously expensive; it is priced for the base case with no cushion.

§1 The Round Trip, in Plain English

Read this if you read nothing else.

For most of a decade Oracle was a boring, profitable, slow-growing software company — the database that runs the back offices of the world, plus a large applications business, throwing off cash and buying back its own stock. Then it became the surprise winner of the AI-infrastructure boom, because it had done something unusual: it had built cheap, standardized data centers and signed enormous multi-year compute contracts with the AI labs, including a reported ~\$300 billion deal with OpenAI. On September 10, 2025, Oracle reported that its **contracted backlog had jumped to \$455 billion**, and the stock did something a company this size almost never does: it rose **36% in a single day**, to a record close of \$324.63, an intraday high of \$345.72.

That was the top. Since then the backlog has grown further — to **\$638 billion** by May 2026 — and the stock has done nothing but fall. It closed at **\$124.21 on July 16, 2026**, down **64% from its peak** and, strikingly, **below its pre-hype level** of \$238 the day before the September blow-out. (The –64% is measured from the intraday high of \$345.72; on a closing basis, \$324.63 to \$124.21, it is –62%.) The backlog nearly doubled; the stock more than halved. The market decided the two facts were not the same thing.

Why? Because a backlog is a promise of *revenue*, and the market spent ten months working out how much *profit* that revenue would carry, who was on the other side of the promises, and what it would cost Oracle to fund the build. The answers, as they emerged: the AI-compute business runs at a **fraction** of Oracle's legacy margins — the one leaked figure put its gross margin near **14%**, versus ~70% for the old business, with an outright **loss on the newest Nvidia chips**. Roughly **half** the backlog is OpenAI, a counterparty whose own finances are, to put it gently, unsettled — enough that S&P **cut Oracle's credit rating to one notch above junk** in July 2026. And to

build the capacity, Oracle's free cash flow went from **+\$11.8 billion to -\$23.7 billion** in two years, funded with debt, while it kept paying a dividend it no longer covers.

So the stock round-tripped. The interesting question is no longer "why did it fall" — it is "is it now cheap, fair, or still expensive?" Our answer: **a touch below fair, with the balance of risk tilted down**. Strip the AI bet out entirely and the old Oracle is worth about **\$72 a share**. At \$124 you are paying roughly **\$52 more** for the AI business — which is neither the giddy number of last September nor nothing, but almost exactly what the AI business is worth if it grows *well but not miraculously*. And when we count the whole company three different ways — by its parts, by its cash flows, by its peers' earnings multiples — the answers cluster at **\$104 to \$120, never above the price**. The method that pays closest attention to the cash Oracle must burn before the build pays off gives the lowest number. There is no discount anywhere in that range for the very real chance the AI business does worse than planned.

§2 Two Oracles in One Ticker

The only way to value this company is to stop averaging its two halves.

Oracle reports as one company, but it is really two businesses with opposite economics, and blending them into a single multiple is how both bulls and bears go wrong. The whole of this note rests on separating them.

- **Oracle A — the annuity.** The database, the applications (Fusion, NetSuite, Health/Cerner), on-premise license support, and the SaaS layer. This is a **~\$49 billion-revenue** business (FY2026) with an operating margin we estimate near **42%**, growing slowly — software license revenue is *flat* (–0.7% in FY26), offset by mid-teens SaaS growth. It is sticky, cash-generative, and about as close to a bond as enterprise software gets. It is most of the profit and almost none of the excitement.
- **Oracle B — the bet.** Oracle Cloud Infrastructure's AI-compute business: renting GPUs and building data centers for the AI labs. **~\$18 billion of revenue** in FY2026 (REPORTED — management's OCI figure on the Q4 call; the 10-K discloses only the combined \$34B "Cloud" line, not an OCI-only breakout), growing ~90%+, but at a margin so thin the company will not print it. This is almost none of the profit and all of the excitement — and all of the \$55.7 billion of capex, the negative free cash flow, and the credit downgrade.

The reason the market can't agree on Oracle is that these two businesses deserve opposite treatment. Oracle A is worth a healthy multiple of stable earnings. Oracle B is worth a *fraction* of its revenue, because its revenue barely carries profit and its assets depreciate on a treadmill (§9). **The \$638 billion backlog is almost entirely Oracle B.** A headline that says "backlog +363%" is describing the low-margin half. Revenue growth, here, dramatically overstates profit growth — and that single sentence is the most important one in this note.

We value the two halves separately, net the debt, and probability-weight three futures for Oracle B (§3–§6) — then check that answer against a consolidated DCF and peer multiples (§7), because a conclusion this dependent on judgment should have to survive methods that don't share its assumptions. Everything below fills in that arithmetic.

§3 Oracle A: What the Annuity Is Worth

The floor under the stock, and it is a real floor.

Legacy Oracle is the part of this business a credit committee or a private-equity buyer would recognize instantly: a wide-moat, high-switching-cost software franchise that has compounded quietly for thirty years.

- **The economics.** Stripping OCI's ~\$18B of AI infrastructure (REPORTED — management's OCI figure) out of the \$67.4B FY2026 total revenue (DISCLOSED) leaves **~\$49.3 billion** across software license and support, SaaS,

hardware, and services. Because the AI business runs near breakeven at the operating line today (\$4), essentially **all** of Oracle's **\$20.6 billion** of FY2026 operating income (DISCLOSED) is attributable to this legacy base — an implied operating margin of **~42%**.

- **The growth.** Modest and durable. On-premise software is in gentle secular decline; SaaS (Fusion, NetSuite, Health) grows mid-teens; support renews at very high rates. Blended, this is a **~5–8% grower** — a software annuity, not a growth stock.
- **The value.** We capitalize the ~\$20.6B of legacy operating income at **15× EBIT** — roughly 18–19× after-tax earnings, a mid-range multiple for a moaty, slow-growing software franchise, and *below* the ~20–25× forward earnings the mega-cap hyperscalers command (MSFT ~20×, AMZN ~26×, GOOGL ~25× forward P/E; REPORTED, July 2026). That yields **~\$309 billion of enterprise value** for Oracle A. Net the \$97.6B of net debt and \$5.0B of preferred, and legacy Oracle is worth about **\$207 billion of equity** — **~\$72 a share**.
- **The sensitivity that matters.** The legacy multiple is a lever: at 17× EBIT (arguably fair for a franchise this durable) legacy is worth ~\$86/share; at 13× (pricing the flat license line as terminal decline) ~\$58. We anchor on 15× and flag the range. Note what the floor is *not* immune to: the \$97.6B of net debt sits *below* the equity, and much of it was raised to fund Oracle B. **If the bet fails, the annuity's owners still owe the bet's debt.** The floor is real but it is a *levered* floor.

Call Oracle A **~\$72 a share**, with a defensible range of ~\$58–86. That is the value you can point to without believing anything about AI. The remaining **~\$52** of the \$124 price is the market's bid for Oracle B.

§4 Oracle B: The Margin Is the Whole Argument

Everything the stock is worth above ~\$72 rests on a number Oracle refuses to disclose.

Oracle B is a real, fast-growing, contracted business. It is also the single hardest thing to value in US large-cap equities right now, because its defining number — the margin it will earn at maturity — is undisclosed, and the only evidence we have points far below what the price implies.

- **What Oracle discloses.** Not much, and what it does disclose points the wrong way. Oracle folded its old revenue lines into "Cloud" and "Software," so there is no OCI-only margin in the 10-K. The one visible fingerprint: the blended **Cloud & Software segment margin fell from 62.8% to 58.9%** year-over-year — a ~390bp compression — *as the low-margin OCI mix grew*. On the Q4 call, management said cloud COGS rose **+56%** against cloud revenue **+47%** (costs outrunning revenue), guided to a **further gross-margin step-down in FY2027**, and framed OCI as growing "into a **30–40% margin profile as it matures**." That last phrase is doing enormous work: it is an **aspiration about the future**, offered precisely because the present does not look like that. [DISCLOSED / REPORTED: Q4 FY26 call]
- **What leaked.** In October 2025, *The Information*, working from internal documents, reported that Oracle's **GPU-cloud deals averaged ~14–16% gross margin, "sometimes below 10%,"** that a quarter of ~\$900M of GPU-rental revenue produced only ~\$125M of gross profit (~14%), and that on the newest Nvidia **Blackwell** systems Oracle ran at a **loss of roughly \$100 million**. This is REPORTED, not DISCLOSED — internal figures Oracle has not confirmed — and we treat it as directional, not audited. But it is the only quantified margin evidence that exists, and it is corroborated in *direction* by the disclosed compression. [REPORTED]
- **The leap the price requires.** Our base case assumes OCI matures to a **22% operating margin**. To get from ~14% *gross* today to 22% *operating* at scale, gross margins must roughly triple and then survive heavy operating costs. That is not impossible — scale economics in cloud are real — but the burden of proof sits squarely on the optimist, and the disclosed trend is running the other way. **The bull case (\$176) assumes a 32% operating margin — more than double the only number anyone outside Oracle has seen.**
- **Why the contract structure caps the upside.** Oracle's management stated on its Q4 call that **\$75 billion** of the large AI contracts are **prepaid or customer-supplied hardware** — the customer either pays up front or brings

its own GPUs (REPORTED; the audited 10-K discloses only ~\$4.6B of significant-financing-component prepayments, so treat the \$75B as management commentary, not a filed line). This genuinely de-risks the *capital* (Oracle is not fronting every dollar of the build). But it also **caps the margin**: hosting a customer's own chips earns colocation economics — thin, in the teens — not cloud economics. The more of OCI that is bring-your-own-hardware, the *harder* the 30–40% aspiration becomes, not easier. The structure that protects the balance sheet also flattens the payoff.

- **The counterparty.** Roughly **half** of the \$638B backlog is OpenAI (per S&P). Our thematic note established that OpenAI's compute commitments are contractual *ceilings* ("up to" amounts), that its funding is partly circular, and that Oracle **has never once named it in an audited filing** — the name appears only in a bond-marketing document. A backlog half-dependent on such a counterparty cannot be valued at face; its realization probability deserves a haircut, which the base and bear cases apply.

The honest summary: Oracle B could be worth a great deal or very little, and the fulcrum is a margin the company will not print. We size the three outcomes next.

§5 What the Price Already Pays For

Invert the question: back out what today's \$124 already assumes, and the picture clarifies.

Rather than argue toward a number, invert the question: at \$124.21, **what is the market already paying for Oracle B?**

- Legacy Oracle is worth **~\$72 a share** (§3).
- The market price is **\$124.21**.
- So the market ascribes **~\$52 a share — about \$151 billion of enterprise value — to Oracle B.**

Now compare that \$151B against our three scenario values for Oracle B (§6): bull ~\$299B, base ~\$145B, bear ~\$24B. **The market's \$151B lands essentially on our base case (\$145B — a ~4% gap, well inside the noise of an exercise built on one undisclosed margin).** This is the cleanest way to state the finding:

> In September 2025, at \$325, the market was paying roughly the **bull-case** value for Oracle B — it was capitalizing the nameplate. After a 64% fall, at \$124, it is paying the **base-case** value. The correction did precisely the work our thematic note wondered whether the equity would do: it moved Oracle from nameplate pricing to base-case pricing.

That is why Oracle is no longer a straightforward short, and why it is also not a bargain. **There is no margin of safety in the price.** You are paying the base-case value for Oracle B, which means you are compensated for neither the well-supported bear case (thin margins, counterparty strain) nor the possibility — real, if the leaked margins are representative — that Oracle B is worth closer to nothing on a standalone economic basis once its depreciation treadmill is charged honestly (§9). To make money from here you need the base case *or better* to be true. The price gives you no discount for the chance it is not.

§6 Three Futures for Oracle B

The scenarios, with the weights argued rather than assumed.

We value Oracle B as a business reaching maturity around FY2030, discounted back. Each scenario is a triplet — revenue, mature operating margin, exit multiple — and the weights carry an explicit, data-driven tilt. (Full arithmetic in the accompanying model; every input traces to the filings.)


- **Bull — \$176 a share (22% weight).** OCI revenue reaches ~\$130B by FY2030 (Catz's stated ramp of \$18→\$32→\$73→\$114→\$144B broadly holds), and the margin matures to **32% operating**, near the low end of management's aspiration. Oracle B is worth ~\$299B. This requires the margin to more than double from the only observed figure — possible if OCI's higher-value platform and database services (not just raw GPU rental) come to dominate the mix, but unproven.
- **Base — \$122 a share (43% weight).** The ramp slips ~30% (to ~\$95B by FY2030) on counterparty attrition, cancelled expansions like the reported Abilene ~600MW, and demand that "frequently changes," and the margin settles at **22% operating** — well above today's economics but below the aspiration. Oracle B is worth ~\$145B. This is roughly what the market pays today.
- **Bear — \$80 a share (35% weight).** Margins stay near the observed ~10% **operating** reality, the capex treadmill (§9) keeps economic returns thin, and OpenAI counterparty strain impairs part of the backlog. OCI reaches only ~\$60B of revenue at a low multiple; Oracle B is worth ~\$24B — and Oracle is left having levered its balance sheet to build a low-return business. The stock falls back toward its levered legacy floor.

Why the weights tilt down. A naive analyst would weight these 25/45/30 and land near \$123 — exactly the current price, which should make one suspicious that "fair" was reverse-engineered. We weight them **22/43/35**, and the reason is evidentiary, not stylistic: three of the four load-bearing facts point the same direction. (1) The base and bull cases imply OCI **gross** margins running **2.5–3× the only observed ~14%** (to reach 22–32% *operating* margins you need gross margins in the mid-30s to mid-40s), and the disclosed trend is compression, not expansion. (2) The prepay/BYOH structure **caps** the achievable margin. (3) The depreciation treadmill means even reported margins **overstate** economic ones (§9). The counterweights are real — legacy is a hard floor, the RPO is contracted, and 64% of the drawdown is behind us — which is why the tilt is *modest*, not severe. But an honest reading of the evidence does not sit at the symmetric midpoint. It sits a little below it. Probability-weighted: **~\$119**.

§7 Cross-Checks: What a DCF and the Comps Say

A single method is a point estimate. Three methods are a triangulation — and where they disagree is where the information is.

The sum-of-the-parts (§3–§6) is one lens. We ran two more, independently, to see whether they corroborate ~\$119 or pull against it. They corroborate the *direction* — Oracle is not cheap — and they pull the *center* modestly lower, for a reason worth understanding.



- **The consolidated DCF says ~\$104, and the reason it is lower than the SOTP is the point.** We built an unlevered free-cash-flow model on the same base-case trajectory (OCI ~\$95B by FY2030), discounted at a

9.25% WACC. It lands at **~\$104 a share** — below the SOTP's \$119 — because a DCF is forced to pay, in full and up front, for the **~\$63 billion of cumulative negative free cash flow Oracle burns in FY2027–FY2029** before the build inflects to cash generation. The SOTP's exit multiple quietly assumes that journey is free; the DCF does not. The DCF is arguably the *more* honest lens on a company spending ~\$90B a year, and it says the SOTP is a touch generous.

- **The DCF's swing factor is the same one as the SOTP's — the treadmill.** Its single most sensitive input is **terminal capex intensity**: how much Oracle must keep spending to replace GPUs in perpetuity. At a normalized 15% of revenue the DCF gives **~\$124**; at a treadmill-like 24% it gives **~\$65**. That \$65-to-\$124 fork is the depreciation question of \$9 restated as a cash-flow assumption — the same fulcrum (how good are the AI-infrastructure economics, really) that drives the SOTP margin. Two independent methods, one crux.
- **Peer multiples say "cheap" — but only the ones that ignore what ails Oracle.** On **forward P/E**, the lens that respects earnings (capex flows through depreciation into EPS), Oracle at 12–20× FY27 EPS is worth **\$90–150, midpoint ~\$120** — corroborating the SOTP. But on **EV/EBITDA** (~\$137) and **EV/Revenue** (~\$119 at 5×, up to \$150 at 6×), Oracle looks cheaper than its megacap peers (which trade ~20–26× forward earnings). **That cheapness is a trap, not a signal:** EV/EBITDA and EV/Revenue flatter Oracle *precisely because they exclude the \$90B of capex and the thin OCI margin* — the two facts that define the bear case. The multiple a bull would cite to call Oracle cheap is the one that looks past the problem. We give those lenses no weight.
- **Where it lands.** The three capex-respecting methods — SOTP **\$119**, DCF **\$104**, P/E **\$120** — span a tight **\$104–120**, averaging **~\$114**. That is our central estimate, ~8% below the \$124.21 price, and every method agrees on the sign: no discount to fair value, with the DCF insisting the tilt is a little more negative than the SOTP alone would suggest. The methods disagree only on *how much* the near-term cash burn should cost — which is, once again, the one thing this entire valuation turns on.

§8 The Balance Sheet Is Now Part of the Thesis

Oracle used to be a cash machine. It is now a leveraged builder, and that changes the equity's risk.

For years Oracle's balance sheet was an afterthought — it borrowed cheaply and bought back stock. The AI build has made the balance sheet central, and not in a good way.

- **Cash generation inverted.** Free cash flow went from **+\$11.8B (FY24) → −\$0.4B (FY25) → −\$23.7B (FY26)** as capex octupled from \$6.9B to \$55.7B. Management guides FY27 capex to **~\$70B net** of customer prepayments (~\$90–95B gross), and S&P projects FY27 **free operating cash flow near −\$42B**. The cash bleed widens before it narrows. [DISCLOSED / REPORTED]
- **Funded with debt, and the coupon is rising.** Total debt reached **\$129.5B** (net debt **\$97.6B**) at May 2026, up from \$92.6B a year earlier. The maturity ladder is comfortably termed out — ~69% falls after FY2031 — but the *new* money is expensive: recent issues carry **4.55–6.89%** coupons (a \$2.75B tranche due 2066 at 6.85%) versus 2–4% legacy paper. As low-coupon debt rolls, interest expense (already **+28% in FY26** to \$4.6B) keeps climbing.
- **The downgrade.** S&P cut Oracle to **BBB− / A-3 on July 9, 2026** — one notch above high yield — naming OpenAI concentration a "key credit risk." Oracle's 5-year CDS is REPORTED at a record ~198bp — wider than its 2008-crisis levels. The company remains investment grade with a genuine cash-generative core, but the *direction* is unambiguous.
- **Still paying a dividend it doesn't cover.** Oracle paid **\$5.8B** of dividends in FY2026 (yield ~1.6%) and repurchased only **\$95M** of stock — buybacks are effectively halted. So it is funding a growing dividend **with debt** while free cash flow is deeply negative. That is a defensible bridge if the AI build inflects to cash generation on schedule, and a slow bleed if it does not.

- **One flattering item to strip.** FY2026 net income of \$17.1B (EPS \$5.83) was lifted by a **+\$3.5B non-operating gain** absent in prior years (likely investment mark-to-market). Operating-income growth (+16.6%) is the clean read; the +37% net-income growth is not repeatable. Normalized, the trailing P/E is higher than the headline 21x.

The balance sheet does not break the thesis, but it removes the old cushion. The equity is now a **residual claim, junior to \$130B of rising-cost debt**, on a business spending ~\$90B a year to build Oracle B. That juniority is why the bear case is worth -36% and not -15%.

§9 The Depreciation Treadmill

Why even the reported margins flatter, and why this ties to our companion note.

Our companion note, *The Depreciation Question*, argues that the hyperscalers' lengthening server lives overstate their current earnings. Oracle is a clean example, and it matters here because it means Oracle B's *true* economic margin is lower than even the thin reported one.

- **Oracle extended its server life from five years to six**, effective the first quarter of FY2025 — a change that mechanically raised FY25 and FY26 margins by spreading depreciation over more years. (There was no further extension in FY26; the life is now six years.) [DISCLOSED: FY26 10-K]
- **The economic life of an AI GPU is widely argued to be three to four years**, as Nvidia ships a new architecture roughly annually and the Blackwell-era parts render prior generations far less competitive for training. A six-year book life against a ~four-year economic life *understates* depreciation and *overstates* operating margin — precisely on the assets at the heart of Oracle B.
- **And the depreciation is about to accelerate hard.** Oracle carries **\$40 billion of construction-in-progress** not yet being depreciated. Reported depreciation already nearly doubled to \$7.6B in FY26; as that \$40B of data-center gear is placed in service, the depreciation line climbs steeply into FY27 and beyond — a headwind to the very margin recovery the bull case requires.

The implication for valuation: the ~14% gross / breakeven-operating economics Oracle B shows *today* are flattered by a six-year life. Charge depreciation on a four-year economic life and the current margin is worse still — which is why our bear case (10% operating at maturity) is not pessimism so much as an honest extrapolation of the disclosed trajectory, and why the aspiration to 30-40% deserves the skepticism we give it.

§10 What Would Change Our View

The falsifiers, stated so the reader can hold us to them.

- **OCI margin disclosure.** The single datum that would move this valuation more than any other is a clean OCI operating margin. If Oracle begins disclosing it and it is **above ~20%**, the base case is validated and the bull gains weight — FV moves toward \$150+. If it prints in the **low teens or below**, the bear case is confirmed and FV moves toward \$90.
- **The FY2027 cash-flow trajectory.** If free operating cash flow lands **materially better than S&P's -\$42B**, the prepayment structure is doing more work than we credit and the balance-sheet risk recedes. If it lands worse, the dividend-on-debt bridge lengthens.
- **The OpenAI relationship.** Any impairment, renegotiation, or cancellation touching the ~half of RPO tied to OpenAI (the reported Abilene ~600MW cut is the template) validates the bear. Conversely, OpenAI's IPO — its S-1 was confidentially submitted June 8, 2026 — could convert the counterparty into an audited, fundable entity and lift the base.

- **The legacy multiple.** Our \$72 floor assumes 15× EBIT for Oracle A. If the market re-rates durable software higher, the floor rises to ~\$86 and the whole valuation shifts up ~\$14 without anything changing at OCI.
- **What would confirm us.** Continued blended-margin compression as OCI scales; the depreciation line climbing faster than revenue; the dividend funded by fresh issuance; the CDS staying wide. Each is on the dashboard.

§11 House View

The call. Oracle is **roughly fairly valued at a central estimate of ~\$114 against a \$124.21 price** — about 8% downside — with the risk distribution **skewed to the downside**. Three independent methods converge on that sign: sum-of-the-parts \$119, a consolidated DCF \$104, peer P/E multiples \$120 (\$7). The equity has already done the hard work: a 64% fall from the September-2025 peak has moved it from pricing the nameplate backlog to pricing our base case. That is why we do not carry the thematic note's bearish prior into a short thesis here — **the correction largely closed the gap the thematic note hypothesized**. But "fairly valued" is not the same as "attractively valued": at \$124 you are paying the base-case value for Oracle B, which offers **no margin of safety** against a bear case that is genuinely well-supported — an observed ~14% gross margin against a price that needs implied gross margins 2.5–3× higher, a depreciation treadmill that flatters even those, and a single counterparty behind half the backlog. **Confidence: moderate.** The legacy floor and the balance-sheet facts are DISCLOSED and firm; the entire spread between our bear and bull — and the gap between our SOTP and our DCF — is one undisclosed number, the OCI margin and its cash-flow cost.

Why the skew is down and not neutral (per our own standard: a tilt must be argued). Three of four load-bearing facts point the same way — the observed-vs-assumed margin gap, the upside cap from prepay/BYOH structures, and the depreciation treadmill — while the counterweights (a real legacy floor, contracted RPO, a drawdown mostly behind us) are strong enough to keep the tilt *modest* rather than turn it into a short. The honest landing is not the symmetric midpoint that happens to equal the share price; it is a little below it, with a fatter left tail.

The steelman against us. The bulls may be right that we are anchoring on a stale leak. OCI is not only GPU rental — its autonomous-database and platform layers carry far higher margins than the 14% rental slice, and as they grow into the mix the blended margin could climb toward the aspiration on its own. The backlog is contracted, FY27 revenue of ~\$90B is reaffirmed, Oracle's cost-engineered data centers are a genuine edge, and the stock has already fallen 64% — most of the bad news is in. If OCI margins surprise to the upside, \$176 is not a ceiling. **We hold our more cautious line because the only observed margin, the disclosed trend, the contract structure, and the depreciation math all point below the price's embedded assumption — and because paying full base-case value for an unproven margin is not a bargain even when it is not a bubble.**

What the analysis implies. Three understandings fall out of the work. First, **Oracle is two companies**, and the ~\$72-a-share annuity is the honest floor beneath the AI bet. Second, **the price already pays for the base case** — the September-to-July correction repriced Oracle from nameplate to base, so the easy asymmetry is gone in both directions. Third, **the entire remaining debate is one undisclosed margin**, and until Oracle prints it, any fair value here — ours included — is a distribution, not a point. We publish the distribution and let the reader weigh it.

§12 Disclosures

Information only. This material is published by TON618 Capital for information purposes only. Nothing here is an offer to sell or a solicitation of an offer to buy any security, fund interest, or digital asset, and nothing here is personalized investment advice or a recommendation regarding any instrument.

Publisher's exclusion. All research is published solely as general, impersonal information of regular circulation. It is not tailored to the objectives or circumstances of any individual and is not issued in connection with compensation from any client. The Fund has no clients and distributes all research free of charge. On that basis it

publishes in reliance on the publisher's exclusion from the definition of "investment adviser" under the Investment Advisers Act of 1940 (§202(a)(11)(D); cf. *Lowe v. SEC*, 472 U.S. 181 (1985)).

Registration & conflicts. TON618 Capital is not registered as an investment adviser or broker-dealer in any capacity. The Fund is a Bitcoin fund and may hold or transact in the securities or digital assets it discusses; material conflicts are disclosed where they exist. As of the date of this note the Fund holds **no position in Oracle or in any other single security named herein**; it maintains a strategic Bitcoin allocation. The Fund receives no compensation from any party in connection with this research.

Use of AI. Artificial intelligence is used in the creation of this research. **A material conflict is disclosed here: this note values a company whose largest AI counterparties compete with, and contract with, the maker of the AI tooling used to produce it.** All methodology and data integrity are reviewed and approved before publication by TON618 Capital's Chief Investment Officer, **Keyth Beck**; errors may nonetheless occur, and readers should verify independently.

CFA. This report was prepared to align with CFA Institute analytical standards (methodology only). CFA® and Chartered Financial Analyst® are registered trademarks owned by CFA Institute. That reference describes the analytical framework applied; it does not imply the report was prepared, reviewed, or authored by a CFA charterholder, and the report is not issued, reviewed, endorsed, certified, or approved by — nor affiliated with — CFA Institute.

Risk & feedback. Past performance is not indicative of future results. Equities are volatile and may result in total loss of capital. A fair-value estimate is not a price target and carries wide uncertainty. Corrections and feedback are welcome — please direct them to CIO **Keyth Beck** at keyth@ton618capital.com.

Source & Verification Note

This note values Oracle from primary filings pulled directly from SEC EDGAR (CIK 0001341439): the **FY2026 10-K** (filed June 22, 2026, accession 0001193125-26-277521) and the **FY2025 10-K** (June 18, 2025), plus the Q4 FY2026 earnings release and call, S&P Global's July 9, 2026 rating action, and market data at the July 16, 2026 close. All three valuation methods — the sum-of-the-parts, the consolidated DCF, and the peer-multiple comps — are reproducible in the two accompanying models; every input traces to a cited figure.

High-confidence, filing-verified (DISCLOSED). Consolidated results FY24/25/26 (revenue \$52.96B → \$57.40B → \$67.36B; operating income \$15.35B → \$17.68B → \$20.61B; diluted GAAP EPS \$3.71 → \$4.34 → \$5.83; ~2,914M diluted shares); the revenue-line split (Cloud \$33.99B, Software \$24.54B, Hardware \$3.08B, Services \$5.74B) and the Cloud & Software segment margin compression (62.8% → 58.9%); RPO \$638B and its 12%/34%/34% recognition schedule; capex \$55.66B, operating cash flow \$31.98B, FCF -\$23.69B; total debt \$129.5B, net debt \$97.6B, the maturity ladder and new-issue coupons, the \$4.95B mandatory-convertible preferred, and the return of total stockholders' equity to positive (\$42.5B); the dividend (\$5.79B paid) and near-zero buyback (\$95M); the six-year server life (extended from five in Q1 FY2025), depreciation \$3.13B → \$3.87B → \$7.62B, and the \$40B of construction-in-progress; and the FY26 non-operating +\$3.55B gain. Market levels — ORCL \$124.21, and the September 10, 2025 record close \$324.63 / intraday high \$345.72 / pre-spike \$238.48 — are from the authoritative Schwab feed.

Deliberately hedged, not asserted as fact (REPORTED / CONSENSUS). The ~14–16% OCI/GPU gross margin and the ~-\$100M Blackwell loss are *The Information's* read of internal documents (~October 2025), **not disclosed in any filing**; we treat them as directional and corroborated only in *direction* by the disclosed segment-margin compression. Management's "30–40% mature OCI margin," the FY27 ~\$90B revenue / ~\$8.05 non-GAAP EPS guide, the Catz OCI ramp (\$18→\$32→\$73→\$114→\$144B), the ~\$18.1B FY26 OCI revenue figure, and the **\$75B prepaid/customer-supplied hardware** total are all management commentary from the Q4 release/call — REPORTED, not audited-10-K line items (the 10-K discloses only the combined \$34B "Cloud" line and ~\$4.6B of

significant-financing prepayments). The S&P downgrade details (BBB-, "key credit risk," ~half of RPO, ~\$42B FY27 FOCF) are consistent across carriers but S&P's own release was not read directly. The ~\$300B Oracle-OpenAI contract and its 2027 start are REPORTED and never confirmed by Oracle in a filing; consistent with the thematic note, **we do not treat any OpenAI figure as disclosed by Oracle**. Sell-side consensus (FY27 revenue ~\$88.8B, EPS ~\$7.49; a ~\$252 median price target we consider stale post-drawdown) is context, not an input to our fair value.

The softest input, named plainly. The entire spread between our bear (\$80) and bull (\$176) is **one undisclosed number — Oracle B's mature operating margin**. Our scenario margins (10% / 22% / 32% operating) are anchored on the disclosed blended-margin compression and the reported ~14% gross figure — reaching them requires implied gross margins the mid-30s to mid-40s, well above the only observed level; they are the load-bearing assumption and the reader should treat them as such. We have argued a modest downward skew from the evidence rather than defaulting to the symmetric midpoint, and we have made the scenario weights explicit so the reader can re-weight them.

Method notes. Fair value triangulates **three independent methods**, all reproducible in the accompanying models. **(1) Sum-of-the-parts (~\$119):** legacy Oracle at 15× EBIT (~\$309B EV), ascribing essentially all consolidated operating income to it on the basis that OCI runs near operating breakeven today (a 17× legacy multiple lifts fair value ~\$14/share); Oracle B valued as a FY2030-maturity business discounted at 10–13% by scenario. **(2) Consolidated unlevered DCF (~\$104):** a six-year FCFF build on the base-case revenue path (FY27 \$90B → FY32 \$178B), blended operating margin dipping then recovering to 30%, D&A rising to ~16% of revenue and capex falling from ~82% to ~24% as the build completes, discounted at a 9.25% WACC with 3.5% terminal growth; the value is highly sensitive to terminal capex intensity (\$65 at 24% of revenue → \$124 at 15%), which is the depreciation-treadmill question in cash-flow form. **(3) Peer multiples (~\$120 on the earnings-respecting lens):** forward P/E of 12–20× on ~\$7.49 FY27 EPS; EV/EBITDA and EV/Revenue are shown but given no weight because they exclude the capex and margin that define the risk. Net debt (\$97.6B) and preferred (\$4.95B) are subtracted from enterprise to equity value throughout; per-share figures use ~2,880M shares. The central estimate (~\$114) is the mean of the three capex-respecting methods; we treat it as the midpoint of a distribution, not a point.

Companion notes. *The Neocloud Fuse* (July 16, 2026) is the thematic parent that designed this test; *The Depreciation Question* supplies the useful-life analysis behind §9. A single-name treatment of **CoreWeave** — the second and contrasting test — follows separately.

Version 1.0 · analyst: TON618 Equity Research. Fair value ≈ \$114 vs \$124.21 (SOTP \$119 / DCF \$104 / P/E \$120); roughly fair, risks skewed down.